

How my visualization tools use little memory: A tale of incrementalization and laziness

Eugene Kirpichov <ekirpichov@gmail.com>

St.-Petersburg Functional Programming User Group

Nov 7, 2012

Outline

Introduction

plot

tplot

Types of stream operations

The transition

Optimization

Lessons learnt

Introduction

<http://jkff.info/software/timeplotters>

Tools for visualizing program behavior from logs, optimized for one-liners.

- ▶ **timeplot** — quantitative graphs about multiple event streams.
 - ▶ Histograms of event durations, of types of events, regular line/dot plots etc.
- ▶ **splot** — Gantt-like chart about a large number of concurrent processes.
 - ▶ Birds-eye view of thread/process/machine interaction patterns

Live demo

Live demo

Why optimize

Because they were slow and a memory hog (boxing, thunks).

- ▶ Took minutes for 100,000's of events
- ▶ Took hours or crashed for 1,000,000's of events
 - ▶ Crashing was the last straw. I couldn't do my job.

How to optimize

Rewrite in C or Java? We can do better!

- ▶ It would be a defeat :)
- ▶ Would rewrite and re-debug everything
- ▶ Would learn nothing

Turned out worth it.

splot

The easy one.

Before

1. Read input into a list
2. Traverse to calibrate axes
3. Traverse to render

After

1. Read input into a list
2. Traverse to calibrate axes
3. Read input into a list
4. Traverse to render

Lazy IO does the rest.

Cost: no output to window, no input from stdin.

Code tour.

tplot: why it can NOT be done

The hard one.

- ▶ Complex data flow: `events:tracks = M:M`.
- ▶ Uses Chart, which keeps data in memory and for good reasons.

Code tour: old version.

tplot: why it CAN be done

Reading too much \neq drawing too much \Rightarrow Chart is not a problem.

Building “data to render” for Chart in 1 pass seemed *possible*.

Code tour: PlotData, Render.hs

tplot: main idea

Push-based:

- ▶ List representation unchanged and hidden
- ▶ Push item
- ▶ Get result (at any moment)

Code tour: StreamSummary.

Live example: average + profiling.

Applicative average.

Types of stream operations

Concept	Logical type	Actual type
Summarize	$[a] \rightarrow b$	Summary a b
Transform	$[a] \rightarrow [b]$	Summary b r \rightarrow Summary a r
Generate	$a \rightarrow [b]$	Summary b r $\rightarrow (a \rightarrow r)$

Composition pipelines become quite funny.

Code tour: RLE etc.

tplot: new architecture

New architecture:

- ▶ Push-based builders for all plot types
- ▶ Driver loop to feed input events to output tracks
- ▶ When done, summarize and render

Code tour: driver loop, plots (vs old code).

tplot: the transition

I wanted to keep things working all the time.

1. Change interface — make the change possible.
2. Change implementation — make the change real.

tplot: the transition

Before: Completely non-incremental.

- ▶ (*interface*) Separate building and rendering
- ▶ Split into modules
- ▶ (*interface*) Explicit 2 passes, both *potentially* incremental
- ▶ Toying with incremental combinators to get a feeling for them
- ▶ (*interface*) Make all plots have an incremental interface (but actually use `toList` bridge)
- ▶ (*implementation*) Incrementalize plots one by one

After: Completely incremental.

Not so easy

Memory leaks due to insufficient strictness:
“push isn't pushing hard enough.”

Question

What and why remains unevaluated until too late?

Profilers didn't help at all. **Debug.Trace** is insufficient: it outputs lines and I need *hierarchy*, not sequence.

Enter Debug.HTrace

```
cabal install htrace
```

“Code” “tour”. Live demo: [average](#).

Why not X?

Iteratee, conduit, pipes, ...

- ▶ Scary, so I tried to get as far as I can without them
- ▶ Ended up very small, simple and nice, so I didn't look back
- ▶ Also educational

Lessons learnt

- ▶ Lazy IO is ok for this task
- ▶ Debugging laziness is hard: I wouldn't start a high-risk real-time project in Haskell
- ▶ Profilers are almost useless for debugging laziness
- ▶ Debug.Trace is better, Debug.HTrace much better
- ▶ Push-based incremental processing is fun and easy